Armonik: An Open-Source Solution for Computation Orchestration and Distribution

Jérôme Gurhem -

Wilfried Kirschenmann - Aneo, Boulogne-Billancourt, France



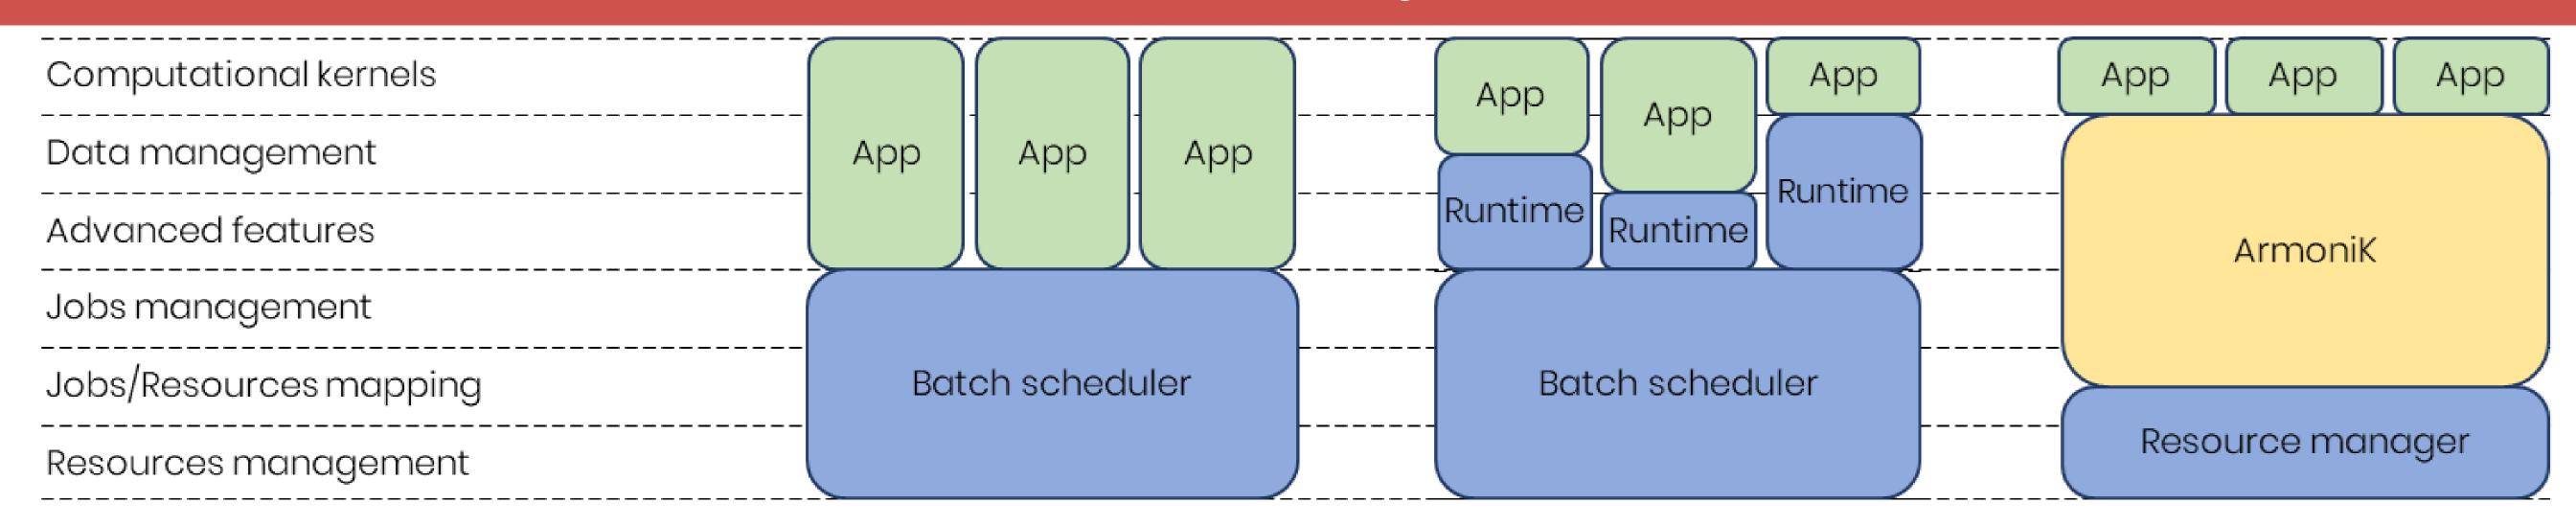
Context

In a world of ever-growing needs for High-Performance Computing (HPC) and massive data processing, **ArmoniK** provides an Open-Source, scalable platform for executing distributed workloads efficiently on heterogeneous infrastructures.

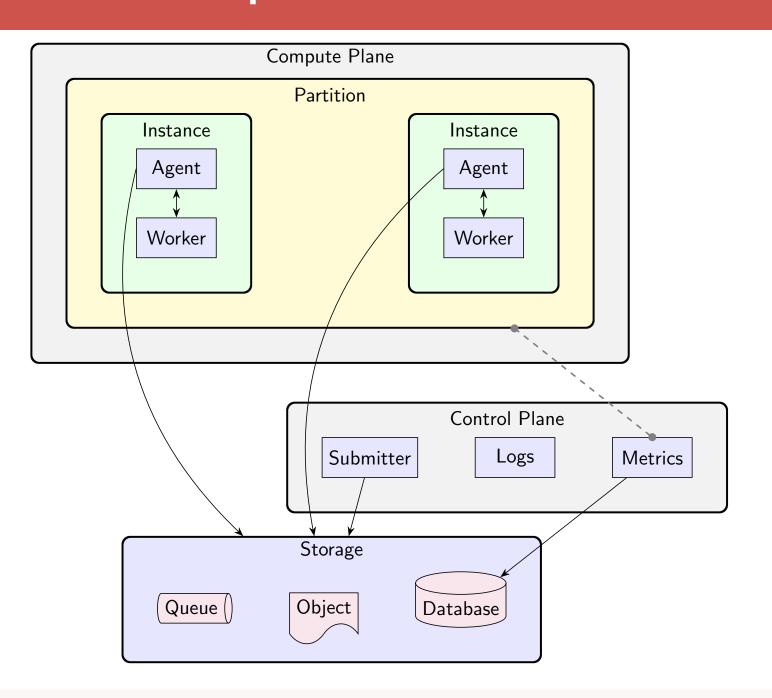
Objectives

- Simplify the development and deployment of distributed computing applications
- Maximize resource utilization across private/public clouds and HPC clusters
- Provide a high-level abstraction for developers

Armonik Positionning in HPC

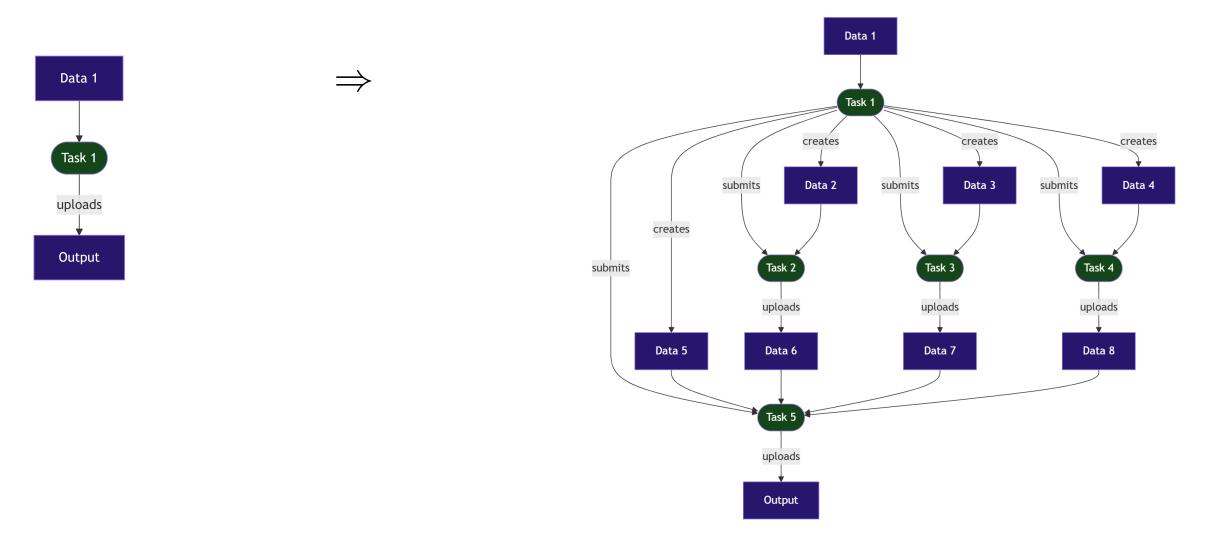


Simplified Architecture



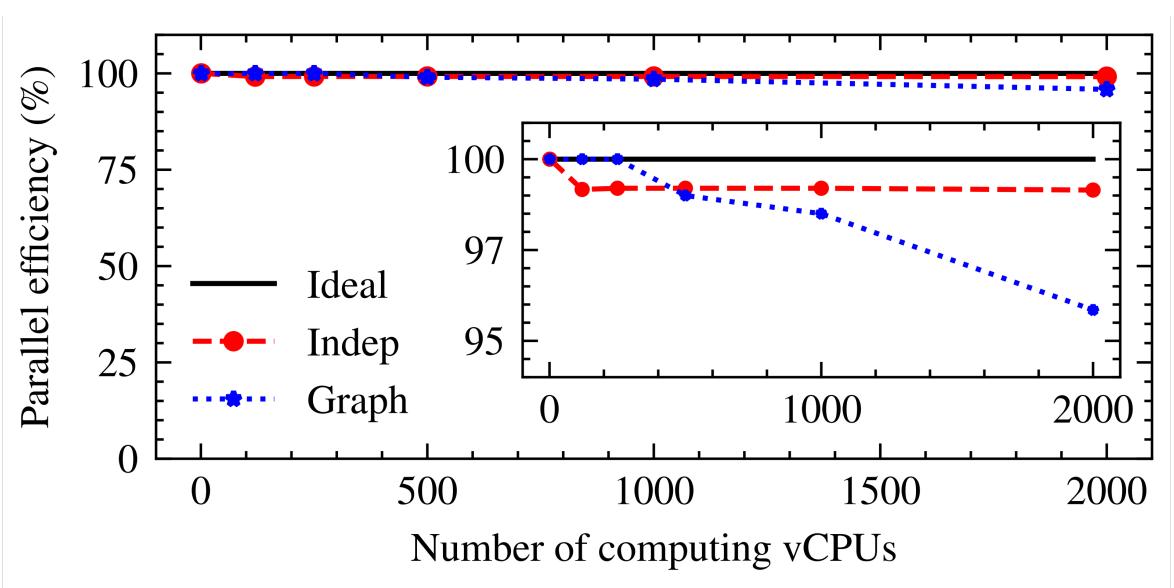
Dynamic Graph

- Dependency graph is not fully known when scheduling starts
- Submissions can happen anytime
- Tasks can submit new tasks
- Tasks can delegate the production of their output to their new tasks



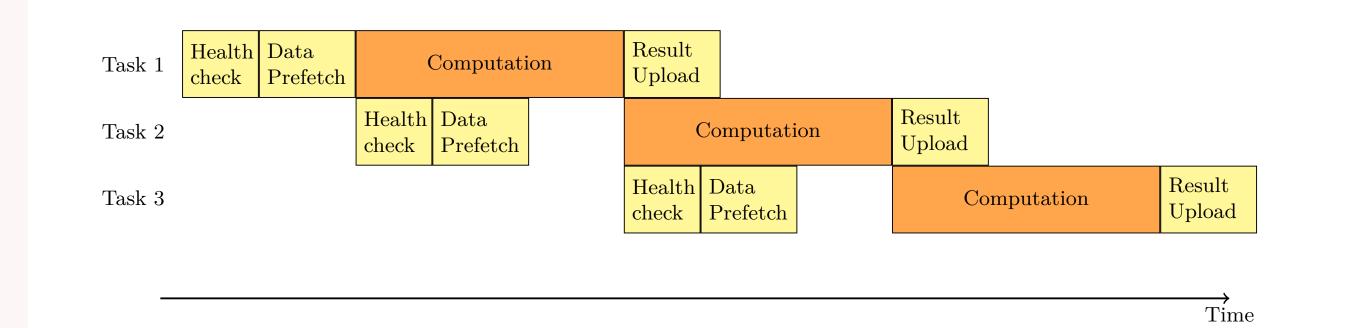
Performance & Scalability

- ► Efficient task retry on failure
- Load-aware scheduling
- Linear scalability on real workloads
- Optimal resource usage on hybrid clusters
- ► Indep: independent tasks workload
- ► Graph : nested fork-join workload



Computations/Comm Overlapping

- ArmoniK is responsible for tasks input and output data management
- ► Allows for automatic communication + scheduling/task execution overlapping
- Automatic Uncoordinated Checkpointing
- Data communications through global storage



Main features

- ► **Observability**: GUIs, CLIs, monitoring APIs, metrics, logs, and traces to understand of the state of the system
- ▶ **Portability**: Easy to transfer an application from one environment to another
- ▷ Officially supported languages: C#, C++, Python, Rust, Java, and JavaScript
- Tasks on different architectures (x86, ARM, GPU, Linux, Windows), applications, environments
- ► Fault Tolerance: Functions without interruption even when one or more nodes fail
 - Allow support for preemptible computing resources
 - Error management at the task level
- Malleability: Support dynamic reconfiguration of the number of allocated resources during execution without interruption
- ► **Resource Sharing**: Allows sharing resources between applications to execute as many as possible at the same
- ► **Modularity**: Modules can be swapped without modifying ArmoniK's code to suit user neeeds and constraints

Use Cases

- Scientific simulation
- Bioinformatics pipelines
- Financial Monte Carlo workloads
- Large-scale data analytics

Conclusion

- ArmoniK simplifies the development of distributed computing applications.
- ► It ensures efficient execution on clouds and HPC clusters through smart orchestration.
- ► Developers benefit from a high-level abstraction and multi-language SDKs.
- Its modular, scalable architecture adapts to changing workloads.
- Integrated observability guarantees reliability and performance.
- ArmoniK enables the next generation of high-performance, data-intensive computing.

https://2025.compas-conference.fr/

© Compas, 26 juin 2025

Bordeaux























